

Unleashing the Power of Deep Reinforcement Learning for Trading

Exploring the Intersection of Reinforcement Learning and Neural Networks

Table of Contents

- [1. Introduction](#)
- [2. Data Engineering](#)
- [3. DRL Environment Architecture](#)
- [4. DRL Agent Architecture](#)

Project Overview

Introduction

Financial markets are a dynamic, complex ecosystem where traditional trading strategies are continually challenged by evolving market conditions, global events, and intricate patterns that define asset prices. In response to this ever-changing landscape, our team has embarked on a pioneering journey to develop a sophisticated DRL-based algorithmic trading system.

Deep Reinforcement Learning, a subset of artificial intelligence, offers the promise of adaptive, self-improving trading strategies. Inspired by the way humans learn through trial and error, our DRL model navigates financial markets, learning and adapting to new information in real-time, all with the goal of maximizing returns while managing risk.

This project represents a convergence of finance, machine learning, and technology. Our DRL model is designed to make informed trading decisions, considering historical data, market indicators, and other relevant factors, ultimately offering the potential to

outperform traditional trading methods. The ability to harness the power of neural networks, paired with reinforcement learning principles, promises a new era in algorithmic trading.

In this project, we will explore the development and implementation of our DRL model, highlighting its capabilities, challenges, and real-world implications in the realm of financial markets. We invite you to journey with us through the intricacies of this innovative approach, where artificial intelligence meets finance, and together, we aim to shape the future of trading.

Objective & Scope

- Collect, clean & analyze Bitcoin crypto market data
- Features Engineering
- Build a DRL trading environment
- Train & improve the agent

Deep Reinforcement Learning Model

- **Reinforcement Learning Algorithm (RLA):** PPO (Proximal Policy Optimization)
- **Neural Network Architecture for Policy (Agent's action):** MLP Policy (Multi-Layer Perceptron Policy)
 - **Architecture Details:**
 - **Input Layer:** The MLP policy takes observations from the environment as input.
 - **Hidden Layers:** The network typically consists of one or more hidden layers, each containing multiple neurons. The specific number of layers and neurons may vary depending on the complexity of the problem.
 - **Activation Functions:** Each layer in the MLP policy may utilize activation functions like ReLU (Rectified Linear Unit) or Tanh.
 - **Output Layer:** The output layer provides the probabilities or values associated with different actions the agent can take.
 - **Training Process:**

- The model is trained using PPO, which involves optimizing the policy network to maximize cumulative rewards while ensuring that policy updates are within a safe margin.

▼ Importing & Downloading Libraries

```
from google.colab import drive  
drive.mount('/content/drive')
```

Mounted at /content/drive

```
!pip install mplfinance
```

Collecting mplfinance

 Downloading mplfinance-0.12.10b0-py3-none-any.whl (75 kB)

 75.0/75.0 kB 1.7 MB/s eta 0:00:00

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from mplfinance) (3.7.1)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from mplfinance) (1.5.3)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfinance)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfinance) (0

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfinanc

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfinanc

Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfinance) (1.

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfinance)

Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfinance) (

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfinance)

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->mplfin

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->mplfinance) (2023.

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib)

Installing collected packages: mplfinance

Successfully installed mplfinance-0.12.10b0



```
!pip install pgmpy  
!pip install pandas-ta
```

```
Collecting pgmpy  
  Downloading pgmpy-0.1.24-py3-none-any.whl (2.0 MB)  
           2.0/2.0 MB 24.1 MB/s eta 0:00:00  
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from pgmpy) (3.2.1)  
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from pgmpy) (1.25.2)  
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from pgmpy) (1.11.4)  
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from pgmpy) (1.2.2)  
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from pgmpy) (1.5.3)  
Requirement already satisfied: pyparsing in /usr/local/lib/python3.10/dist-packages (from pgmpy) (3.1.1)  
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (from pgmpy) (2.1.0+cu121)  
Requirement already satisfied: statsmodels in /usr/local/lib/python3.10/dist-packages (from pgmpy) (0.14.1)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from pgmpy) (4.66.2)  
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from pgmpy) (1.3.2)  
Requirement already satisfied: opt-einsum in /usr/local/lib/python3.10/dist-packages (from pgmpy) (3.3.0)  
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->pgmpy) (2023.4)  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->pgmpy) (2023.4)  
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->pgmpy) (2023.6.0)  
Requirement already satisfied: patsy>=0.5.4 in /usr/local/lib/python3.10/dist-packages (from statsmodels->pgmpy) (0.5.6)  
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.10/dist-packages (from statsmodels->pgmpy) (23.1.3)  
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch->pgmpy) (3.13.1)  
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from torch->pgmpy) (4.9.0)  
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch->pgmpy) (1.12)  
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch->pgmpy) (3.1.3)  
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch->pgmpy) (2023.6.0)  
Requirement already satisfied: triton==2.1.0 in /usr/local/lib/python3.10/dist-packages (from torch->pgmpy) (2.1.0)  
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy>=0.5.4->statsmodels->pgmpy) (1.16.0)  
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch->pgmpy) (2023.4)  
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch->pgmpy) (1.3.0)  
Installing collected packages: pgmpy  
Successfully installed pgmpy-0.1.24  
Collecting pandas-ta  
  Downloading pandas_ta-0.3.14b.tar.gz (115 kB)  
           115.1/115.1 kB 7.9 MB/s eta 0:00:00  
    Preparing metadata (setup.py) ... done  
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from pandas-ta) (1.5.3)  
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->pandas-ta) (2023.4)  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->pandas-ta) (2023.4)  
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.10/dist-packages (from pandas->pandas-ta) (1.25.3)
```

```
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas)
Building wheels for collected packages: pandas-ta
  Building wheel for pandas-ta (setup.py) ... done
    Created wheel for pandas-ta: filename=pandas_ta-0.3.14b0-py3-none-any.whl size=218907 sha256=91d87a0f0437ad983a4cccd5b
      Stored in directory: /root/.cache/pip/wheels/69/00/ac/f7fa862c34b0e2ef320175100c233377b4c558944f12474cf0
Successfully built pandas-ta
Installing collected packages: pandas-ta
Successfully installed pandas-ta-0.3.14b0
```

```
!pip install stable-baselines3
```

```
Collecting stable-baselines3
  Downloading stable_baselines3-2.2.1-py3-none-any.whl (181 kB)
    ━━━━━━━━━━━━━━━━ 181.7/181.7 kB 3.0 MB/s eta 0:00:00
Collecting gymnasium<0.30,>=0.28.1 (from stable-baselines3)
  Downloading gymnasium-0.29.1-py3-none-any.whl (953 kB)
    ━━━━━━━━━━━━━━ 953.9/953.9 kB 35.3 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from stable-baselines3) (1.25.2)
Requirement already satisfied: torch>=1.13 in /usr/local/lib/python3.10/dist-packages (from stable-baselines3) (2.1.0+cu113)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-packages (from stable-baselines3) (2.2.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from stable-baselines3) (1.5.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from stable-baselines3) (3.7.1)
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/python3.10/dist-packages (from gymnasium<0.30,>=0.28.1->stable-baselines3)
Collecting farama-notifications>=0.0.1 (from gymnasium<0.30,>=0.28.1->stable-baselines3)
  Downloading Farama_Notifications-0.0.4-py3-none-any.whl (2.5 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.13->stable-baselines3) (3.5.0)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.13->stable-baselines3) (1.10.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.13->stable-baselines3) (2.6.3)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.13->stable-baselines3) (3.1.2)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch>=1.13->stable-baselines3) (2.2.0)
Requirement already satisfied: triton==2.1.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.13->stable-baselines3) (2.1.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->stable-baselines3) (1.0.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->stable-baselines3) (0.10.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->stable-baselines3) (4.22.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->stable-baselines3) (1.0.1)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->stable-baselines3) (20.4)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->stable-baselines3) (8.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->stable-baselines3) (2.4.7)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->stable-baselines3) (2.8.1)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->stable-baselines3) (2020.1)
```

```
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.13->stable)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch>=1.13->stable)
Installing collected packages: farama-notifications, gymnasium, stable-baselines3
Successfully installed farama-notifications-0.0.4 gymnasium-0.29.1 stable-baselines3-2.2.1
```

```
!pip install 'shimmy>=0.2.1'
```

```
Collecting shimmy>=0.2.1
  Downloading Shimmy-1.3.0-py3-none-any.whl (37 kB)
Requirement already satisfied: numpy>=1.18.0 in /usr/local/lib/python3.10/dist-packages (from shimmy>=0.2.1) (1.25.2)
Requirement already satisfied: gymnasium>=0.27.0 in /usr/local/lib/python3.10/dist-packages (from shimmy>=0.2.1) (0.29.0)
Requirement already satisfied:云cloudpickle>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from gymnasium>=0.27.0->stable)
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/python3.10/dist-packages (from gymnasium>=0.27.0->stable)
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/lib/python3.10/dist-packages (from gymnasium>=0.27.0->stable)
Installing collected packages: shimmy
Successfully installed shimmy-1.3.0
```



```
import pandas as pd
import gym
import numpy as np
from stable_baselines3 import PPO
from stable_baselines3.common.vec_env import DummyVecEnv
import warnings
warnings.filterwarnings("ignore")
```

```
import numpy as np
import pandas as pd
from pandas.plotting import scatter_matrix
# from pandas_ta import wma, macd, stoch, adx, atr, hma, kama, cmo, zscore, qstick

import calendar
import os
import random
import itertools

from matplotlib.colors import TwoSlopeNorm, LinearSegmentedColormap
import matplotlib.pyplot as plt
import matplotlib.colors as mcolors
import matplotlib.cm as cm
from matplotlib.cm import ScalarMappable
from matplotlib.colors import Normalize
import seaborn as sns
import plotly.express as px

from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
import statsmodels.api as sm
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.tsa.seasonal import STL
# from statsmodels.stats.outliers_influence import variance_inflation_factor

from scipy.stats import pearsonr
from scipy import stats
from scipy.stats.mstats import winsorize
from scipy.signal import periodogram

import warnings
warnings.filterwarnings('ignore')
warnings.simplefilter('ignore')

import networkx as nx
```

▼ Data Engineering

```
df = pd.read_csv('/content/drive/MyDrive/IS2/crypto_data.csv')
```

```
df.head()
```

	date	symbol	open	high	low	close	volume	usdt	tradecount	token	hour	day
0	2020-12-25 05:00:00	1INCHUSDT	0.2000	3.0885	0.2000	2.5826	35530516	48768	1INCH	5	Friday	
1	2020-12-25 06:00:00	1INCHUSDT	2.5824	2.6900	2.2249	2.5059	22440875	31099	1INCH	6	Friday	
2	2020-12-25 07:00:00	1INCHUSDT	2.5152	2.8870	2.3609	2.6237	21300426	33001	1INCH	7	Friday	
3	2020-12-25 08:00:00	1INCHUSDT	2.6318	2.8247	2.4650	2.6134	17491813	30459	1INCH	8	Friday	
4	2020-12-25 09:00:00	1INCHUSDT	2.6104	2.7498	2.5629	2.6365	9919400	21023	1INCH	9	Friday	

```
df.shape
```

```
(1120061, 11)
```

```
df.columns
```

```
Index(['date', 'symbol', 'open', 'high', 'low', 'close', 'volume usdt',
       'tradecount', 'token', 'hour', 'day'],
      dtype='object')
```

```
start_date = '2020-08-17 04:00:00'
end_date = '2023-10-19 23:00:00'
```

```
data_df = df.copy(deep=True)
```

```

data_df = data_df[(data_df['token'] == 'BTC') & (data_df['date'] >= start_date) & (data_df['date'] <= end_date)]

data_df.columns

Index(['date', 'open', 'high', 'low', 'close', 'volume usdt', 'tradecount',
       'hour', 'day', 'ema_13', 'ema_25', 'ema_32', 'ema_100', 'ema_200',
       'vol_close', 'vol_close_ema_3', 'vol_close_ema_6', 'vol_close_ema_12'],
      dtype='object')

```

▼ Mapping 'day' days to numerical values

```

day_mapping = {
    'Monday': 1,
    'Tuesday': 2,
    'Wednesday': 3,
    'Thursday': 4,
    'Friday': 5,
    'Saturday': 6,
    'Sunday': 7
}

data_df['day'] = data_df['day'].apply(lambda x: day_mapping[x])

first_appearance = df.groupby('token')['date'].min()
token_counts = df['token'].value_counts()
token_info = pd.DataFrame({'Count': token_counts, 'First Appearance': first_appearance})
token_info['Number of Days'] = (token_info['Count'] / 24).astype(int)
token_info = token_info.sort_values(by='Number of Days', ascending=False)

print(token_info)

      Count First Appearance Number of Days
ETH    54116 2017-08-17 04:00:00          2254
BTC    54116 2017-08-17 04:00:00          2254
BNB    52032 2017-11-06 03:00:00          2168

```

LTC	51285	2017-12-13 03:00:00	2136
ADA	48198	2018-04-17 04:00:00	2008
ETC	46942	2018-06-12 02:00:00	1955
TRX	46905	2018-06-11 11:00:00	1954
LINK	41665	2019-01-16 10:00:00	1736
XMR	40316	2019-03-15 04:00:00	1679
MATIC	39243	2019-04-26 15:00:00	1635
ATOM	39182	2019-04-29 04:00:00	1632
ALGO	37900	2019-06-22 00:00:00	1579
DOGE	37576	2019-07-05 12:00:00	1565
DENT	36320	2019-08-27 04:00:00	1513
RVN	35616	2019-09-25 12:00:00	1484
HBAR	35528	2019-09-29 04:00:00	1480
BCH	34086	2019-11-28 10:00:00	1420
COMP	29059	2020-06-25 06:00:00	1210
MKR	28382	2020-07-23 14:00:00	1182
SOL	27934	2020-08-11 06:00:00	1163
BAL	27934	2020-08-11 06:00:00	1163
CRV	27840	2020-08-15 04:00:00	1160
DOT	27749	2020-08-18 23:00:00	1156
SUSHI	27445	2020-09-01 11:00:00	1143
UNI	27069	2020-09-17 03:00:00	1127
AVAX	26926	2020-09-22 06:00:00	1121
AAVE	26373	2020-10-15 03:00:00	1098
FIL	26363	2020-10-15 17:00:00	1098
1INCH	24672	2020-12-25 05:00:00	1028
SHIB	21413	2021-05-10 11:00:00	892
ICP	21399	2021-05-11 01:00:00	891
DYDX	18477	2021-09-09 02:00:00	769

```
unique_tokens = df['token'].unique()
```

```
unique_tokens
```

```
array(['1INCH', 'AAVE', 'ADA', 'ALGO', 'ATOM', 'AVAX', 'BAL', 'BCH',
       'BNB', 'BTC', 'COMP', 'CRV', 'DENT', 'DOGE', 'DOT', 'DYDX', 'ETC',
       'ETH', 'FIL', 'HBAR', 'ICP', 'LINK', 'LTC', 'MATIC', 'MKR', 'RVN',
       'SHIB', 'SOL', 'SUSHI', 'TRX', 'UNI', 'XMR'], dtype=object)
```

EDA

```
data = pd.read_csv('/content/drive/MyDrive/IS2/crypto_data.csv')

data['date'] = pd.to_datetime(data['date'], format='%Y-%m-%d %H:%M:%S')
data['month'] = data['date'].dt.month
data['week'] = data['date'].dt.isocalendar().week
data['week'] = data['week'].astype(int)
data['year'] = data['date'].dt.year

# new feature 'volatility'(%)
data['volatility'] = (data['high'] - data['low']) / data['low'] * 100
data['hc_change'] = (data['close'] - data['high']) / data['high'] * 100
data['lc_change'] = (data['close'] - data['low']) / data['low'] * 100

#new feature 'price_change_percent' (change in % from close to close)

data = data.sort_values(by=['token', 'date'])
data['price_change_percent'] = (data.groupby('token')['close'].pct_change() * 100).round(2)

data.head()
```

	date	symbol	open	high	low	close	volume usdt	tradecount	token	hour	day	month	week	year	vola
0	2020-12-25 05:00:00	1INCHUSDT	0.2000	3.0885	0.2000	2.5826	35530516	48768	1INCH	5	Friday	12	52	2020	1444.1
1	2020-12-25 06:00:00	1INCHUSDT	2.5824	2.6900	2.2249	2.5059	22440875	31099	1INCH	6	Friday	12	52	2020	20.1

```
#renaming column

data.rename(columns={'volume usdt': 'volume'}, inplace=True)
data.shape

(1120061, 18)

nan_values = data.isna()
any_missing_values = nan_values.any().any()

any_missing_values_in_column = nan_values.any()

missing_value_count = data.isnull().sum()

for column, has_missing in any_missing_values_in_column.items():
    if has_missing:
        count = missing_value_count[column]
        print(f"----> Column '{column}' has {count} missing values.")

print("\nMissing Values in the Entire DataFrame?")
print(any_missing_values)

print("\nMissing Values in Each Column?")
print(any_missing_values_in_column)

print("\nMissing Value Counts in Each Column:")
print(missing_value_count)

----> Column 'price_change_percent' has 32 missing values.

Missing Values in the Entire DataFrame?
True

Missing Values in Each Column?
date                  False
symbol                False
open                  False
high                 False
low                  False
```

```
close           False
volume          False
tradecount      False
token           False
hour            False
day             False
month           False
week            False
year            False
volatility      False
hc_change       False
lc_change       False
price_change_percent True
dtype: bool
```

Missing Value Counts in Each Column:

```
date            0
symbol          0
open            0
high            0
low             0
close           0
volume          0
tradecount      0
token           0
hour            0
day             0
month           0
week            0
year            0
volatility      0
hc_change       0
lc_change       0
price_change_percent 32
dtype: int64
```

#We are expecting 32 unique tokens. Which explains the 32 missing values for price_change

```
un_token = data['token'].nunique()
un_token
```

```
#The first row of each token has a NaN value for price_change'. Normal because there was no input before since it's the fir
```

```
data['price_change_percent'] = data.groupby('token')['price_change_percent'].fillna(0)
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1120061 entries, 0 to 1120060
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   date             1120061 non-null   datetime64[ns]
 1   symbol            1120061 non-null   object  
 2   open              1120061 non-null   float64 
 3   high              1120061 non-null   float64 
 4   low               1120061 non-null   float64 
 5   close              1120061 non-null   float64 
 6   volume             1120061 non-null   int64  
 7   tradecount         1120061 non-null   int64  
 8   token              1120061 non-null   object  
 9   hour               1120061 non-null   int64  
 10  day                1120061 non-null   object  
 11  month              1120061 non-null   int64  
 12  week               1120061 non-null   int64  
 13  year               1120061 non-null   int64  
 14  volatility          1120061 non-null   float64 
 15  hc_change           1120061 non-null   float64 
 16  lc_change           1120061 non-null   float64 
 17  price_change_percent 1120061 non-null   float64 
dtypes: datetime64[ns](1), float64(8), int64(6), object(3)
memory usage: 162.4+ MB
```

```
data_df = data.copy(deep=True)
df = data.copy(deep=True)
df2 = data.copy(deep=True)
df3 = data.copy(deep=True)
```

```
#Checking all unique tokens in the dataset

u_token = df['token'].unique()
u_token

array(['1INCH', 'AAVE', 'ADA', 'ALGO', 'ATOM', 'AVAX', 'BAL', 'BCH',
       'BNB', 'BTC', 'COMP', 'CRV', 'DENT', 'DOGE', 'DOT', 'DYDX', 'ETC',
       'ETH', 'FIL', 'HBAR', 'ICP', 'LINK', 'LTC', 'MATIC', 'MKR', 'RVN',
       'SHIB', 'SOL', 'SUSHI', 'TRX', 'UNI', 'XMR'], dtype=object)

t_counts = df['token'].value_counts()
tc_df = t_counts.reset_index()
tc_df.columns = ['token', 'count']
tc_df = tc_df.sort_values(by='count', ascending=False)

d_counts = df['day'].value_counts()
d_df = d_counts.reset_index()
d_df.columns = ['day', 'count']
d_df = d_df.sort_values(by='count', ascending=False)

h_counts = df['hour'].value_counts()
h_df = h_counts.reset_index()
h_df.columns = ['hour', 'count']
h_df = h_df.sort_values(by='count', ascending=False)

m_counts = df['month'].value_counts()
m_df = m_counts.reset_index()
m_df.columns = ['Month', 'count']
m_df = m_df.sort_values(by='count', ascending=False)
```

```
fig, axs = plt.subplots(4, 1, figsize=(16, 12))

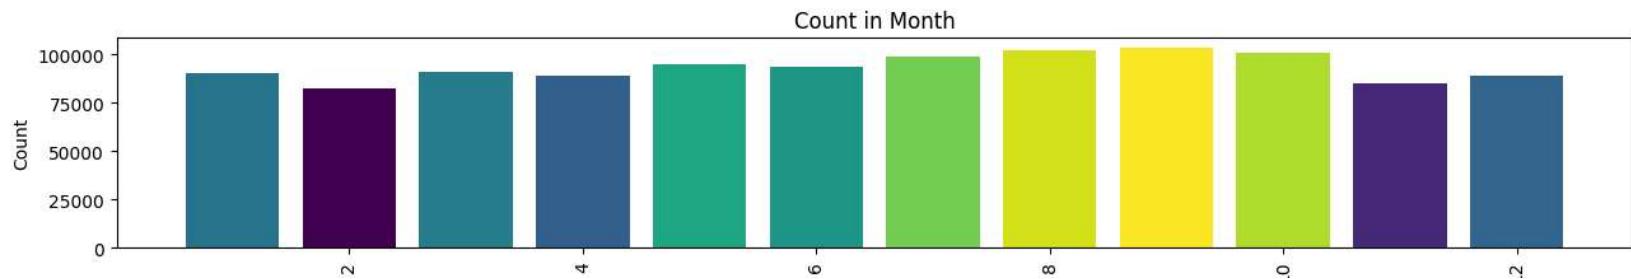
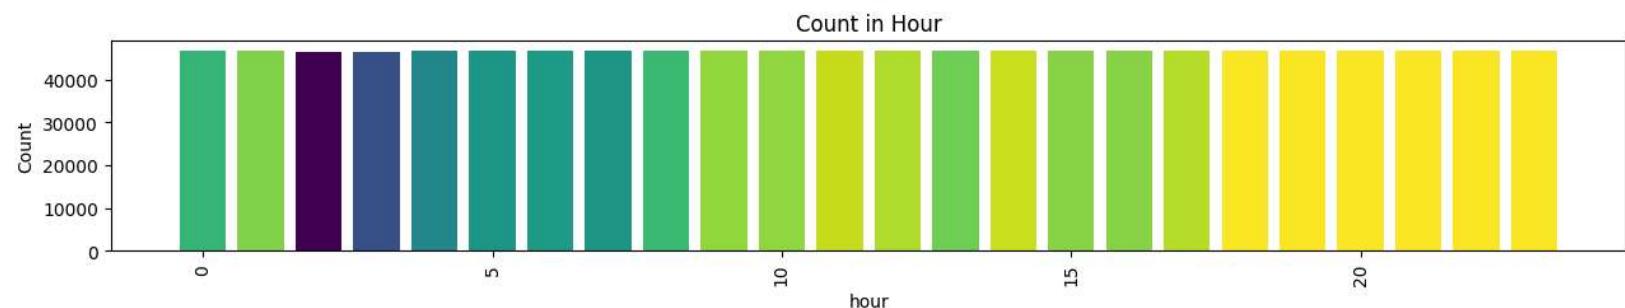
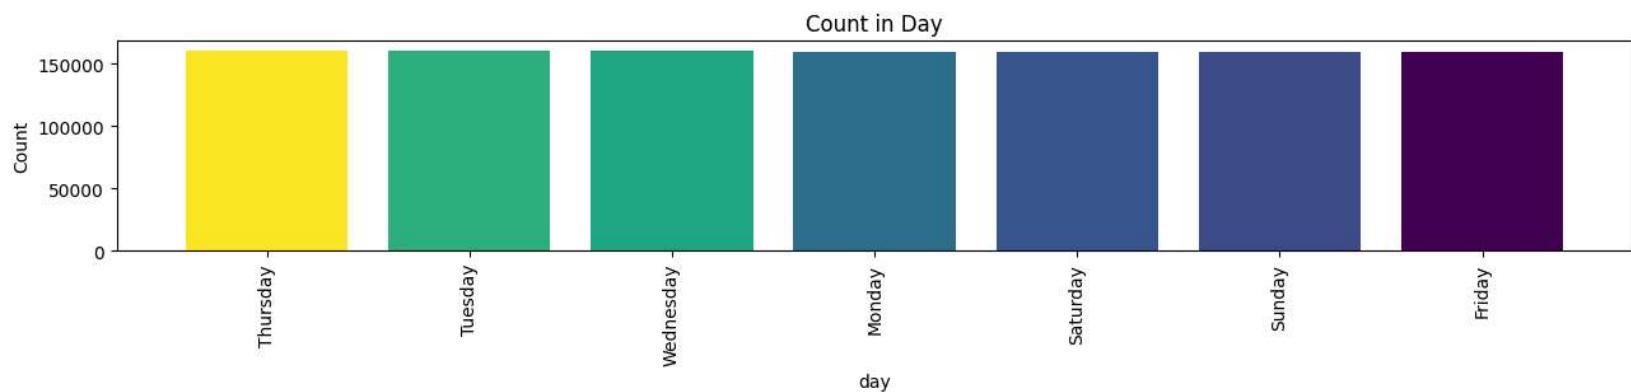
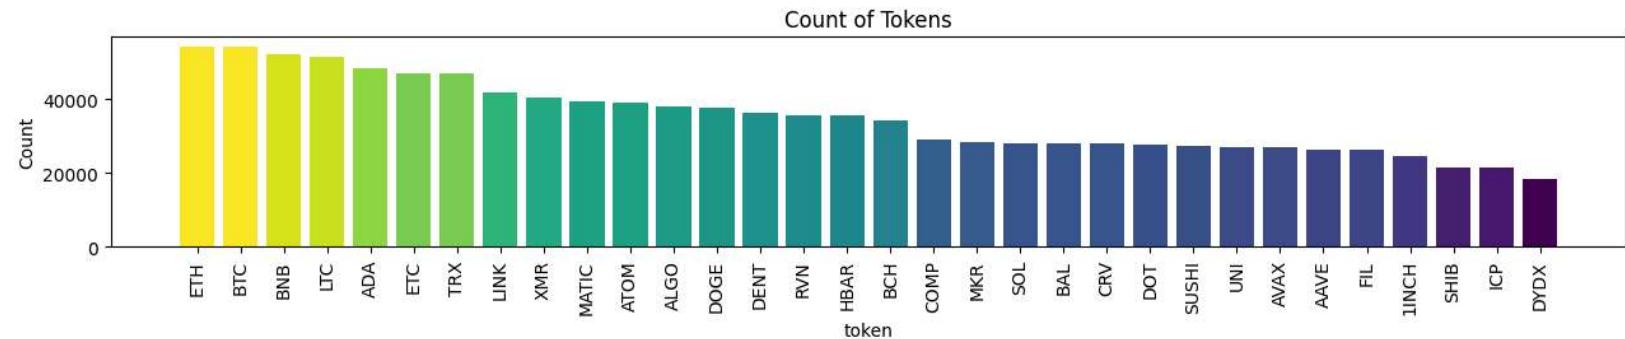
for i, (df, title) in enumerate([(tc_df, 'Count of Tokens'),
                                 (d_df, 'Count in Day'),
                                 (h_df, 'Count in Hour'),
                                 (m_df, 'Count in Month')]):
    colormap = plt.cm.get_cmap('viridis')
    norm = Normalize(vmin=df['count'].min(), vmax=df['count'].max())

    bars = axs[i].bar(df[df.columns[0]], df['count'], color=colormap(norm(df['count'])))

    axs[i].set_xlabel(df.columns[0])
    axs[i].set_ylabel('Count')
    axs[i].set_title(title)
    axs[i].tick_params(axis='x', rotation=90)

    sm = ScalarMappable(cmap=colormap, norm=norm)
    sm.set_array([])
    cbar = plt.colorbar(sm, ax=axs[i])

plt.tight_layout()
plt.show()
```



00:00 00:00 00:00 00:00 Month

```
df2['symbol'] = label_encoder.fit_transform(df2['symbol'])

day_mapping = {
    'Monday': 0,
    'Tuesday': 1,
    'Wednesday': 2,
    'Thursday': 3,
    'Friday': 4,
    'Saturday': 5,
    'Sunday': 6
}

df2['day'] = df2['day'].map(day_mapping)

scat_filt = ['hour', 'day', 'week', 'month', 'volatility', 'tradecount', 'volume']

#btc_scat = df2[(df2['token'] == 'BTC') & (df2['year'] == 2020)]
btc_scat = df2[(df2['token'] == 'BTC')]

btc_scat = btc_scat[scat_filt]
```

```
btc_scat.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 54116 entries, 317303 to 371418
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   hour        54116 non-null   int64  
 1   day         54116 non-null   int64  
 2   week        54116 non-null   int64  
 3   month       54116 non-null   int64  
 4   volatility  54116 non-null   float64 
 5   tradecount  54116 non-null   int64  
 6   volume      54116 non-null   int64  
dtypes: float64(1), int64(6)
memory usage: 3.3 MB
```

```
percentile_threshold = 99.5
percentile_value = np.percentile(btc_scat['volatility'], percentile_threshold)
btc_scat = btc_scat[btc_scat['volatility'] <= percentile_value]
```

```
scatter_columns = ['hour', 'day', 'week', 'month']

num_rows = len(scatter_columns) // 2 + (len(scatter_columns) % 2 > 0)

fig, axes = plt.subplots(num_rows, 2, figsize=(8 * 2, 6 * num_rows))

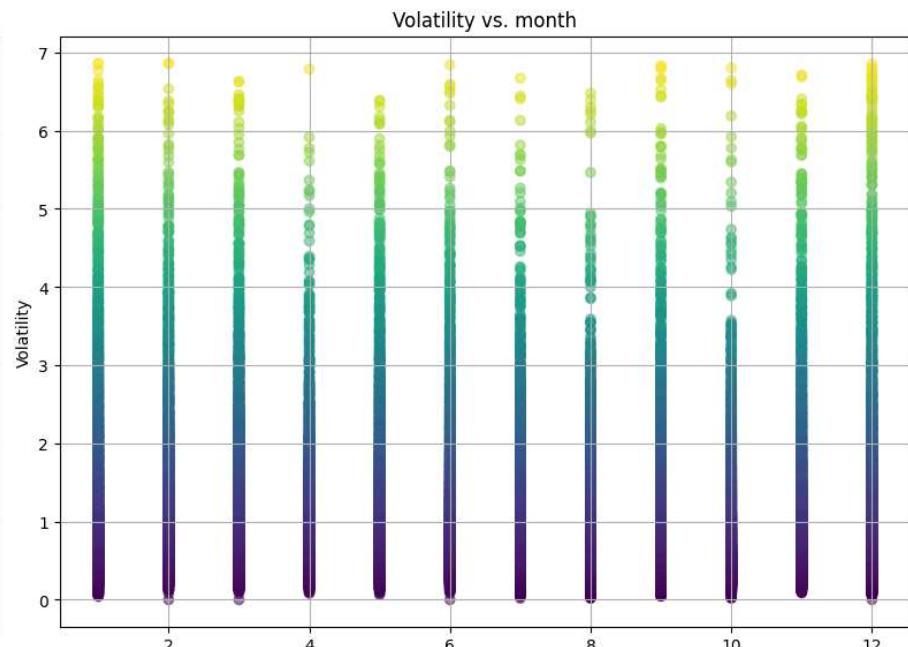
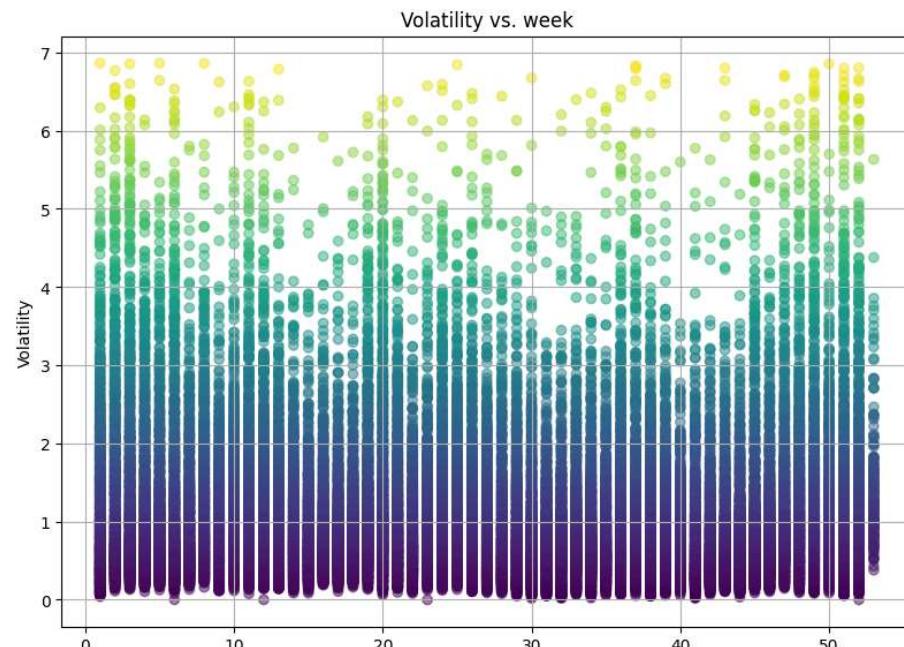
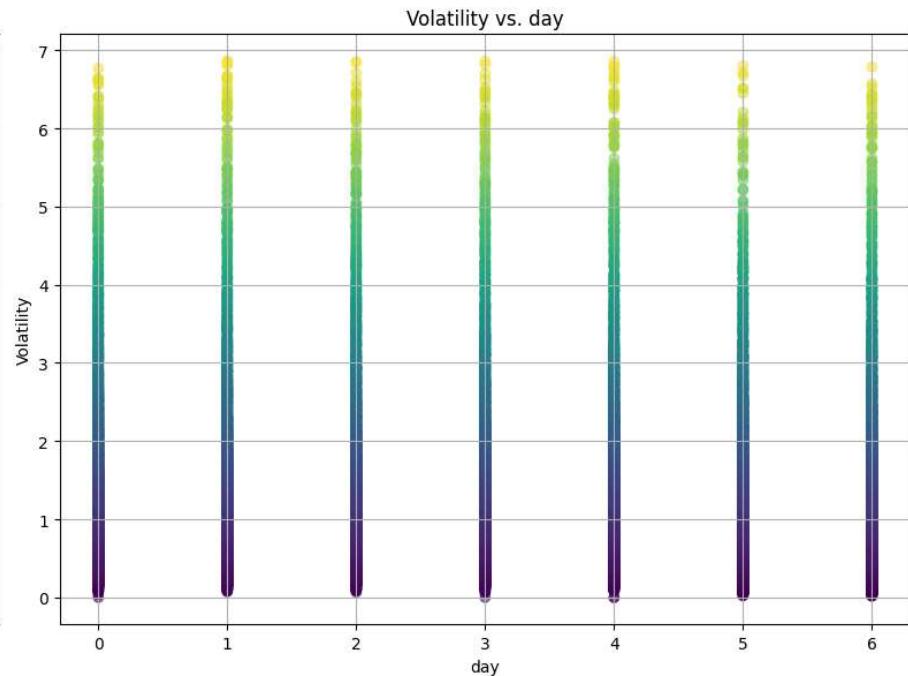
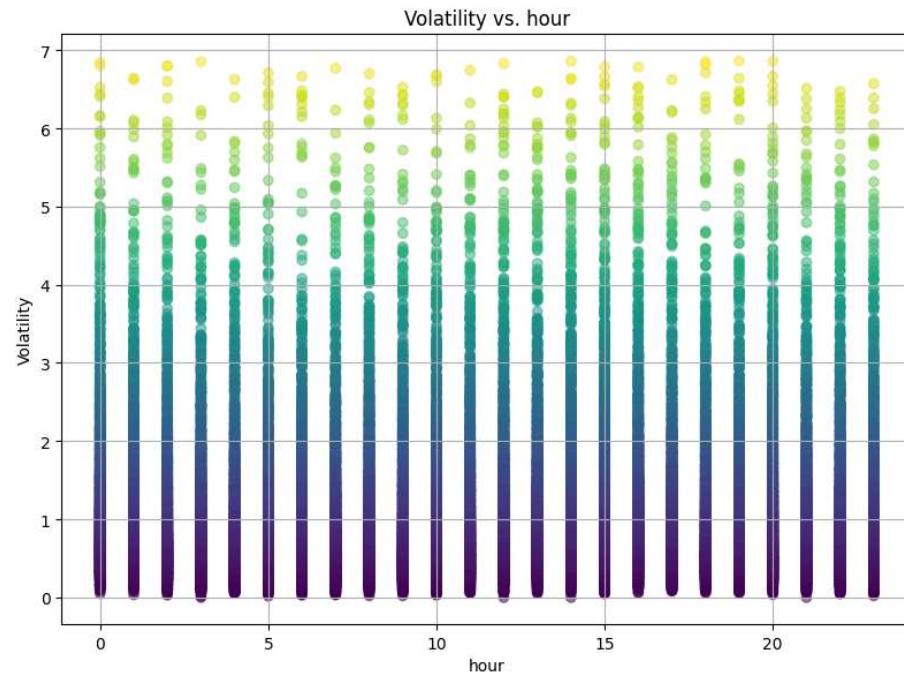
for i, column in enumerate(scatter_columns):
    plt.subplot(num_rows, 2, i + 1)

    colors = btc_scat['volatility']

    plt.scatter(btc_scat[column], btc_scat['volatility'], c=colors, cmap=cm.viridis, alpha=0.5)
    plt.xlabel(column)
    plt.ylabel('Volatility')
    plt.title(f"Volatility vs. {column}")
    plt.legend().set_visible(False)
    plt.grid(True)

plt.tight_layout()
plt.show()
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an u
WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an u
WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an u
WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an u



week

month

```
btc_df = df2[df2['token'] == 'BTC']
```

```
btc_df_2023_october = btc_df[(btc_df['date'].dt.year == 2023) & (btc_df['date'].dt.month == 10)]  
  
dates = btc_df_2023_october['date']  
open_prices = btc_df_2023_october['open']  
high_prices = btc_df_2023_october['high']  
low_prices = btc_df_2023_october['low']  
  
plt.figure(figsize=(16, 8))  
plt.gca().set_facecolor('black')  
  
plt.plot(dates, high_prices, label='High', color='green')  
plt.plot(dates, low_prices, label='Low', color='red')  
  
plt.fill_between(dates, low_prices, high_prices, where=(high_prices > low_prices), alpha=0.3, color='white')  
  
plt.xlabel('Date')  
plt.ylabel('Price')  
plt.title('BTC Price Range (High-Low) for October 2023')  
plt.legend()  
  
plt.grid(color='white', alpha=0.2)  
plt.show()
```

BTC Price Range (High-Low) for October 2023



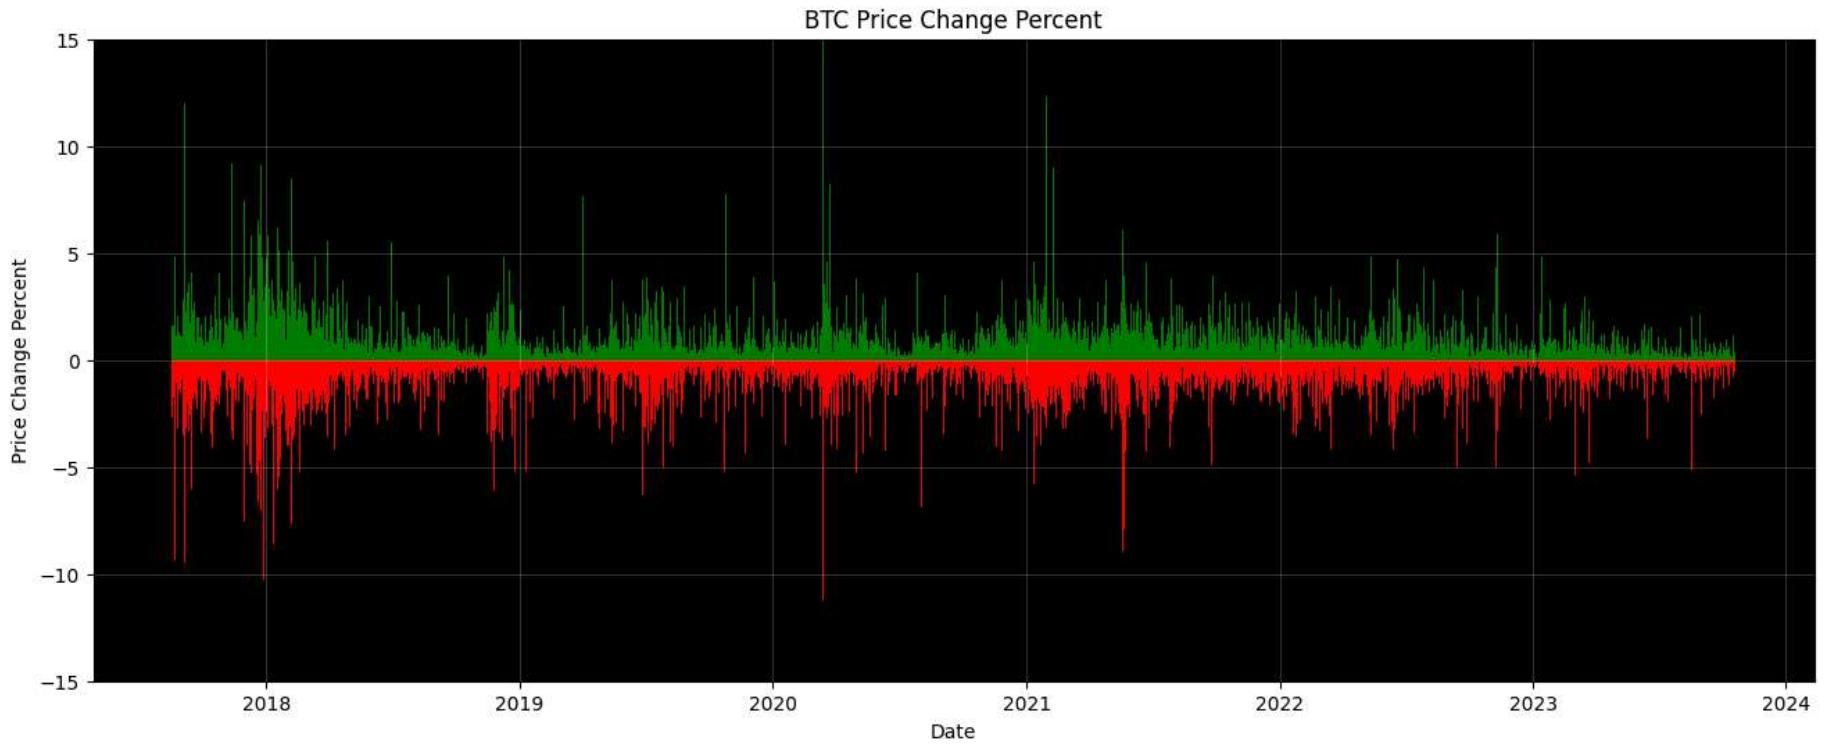
```
price_change_percent = btc_df['price_change_percent']

plt.figure(figsize=(16, 6))
plt.gca().set_facecolor('black')
bars = plt.bar(btc_df['date'], price_change_percent, color=['g' if p >= 0 else 'r' for p in price_change_percent])
plt.ylim(-15, 15)

plt.xlabel('Date')
plt.ylabel('Price Change Percent')
plt.title('BTC Price Change Percent')

plt.grid(color='white', alpha=0.2)

plt.show()
```



- ▼ Creating new columns EMAs (EMA 13 25 32 100 200) - For Moving Average

```
data_df['ema_13'] = data_df['close'].ewm(span=13).mean()
data_df['ema_25'] = data_df['close'].ewm(span=25).mean()
data_df['ema_32'] = data_df['close'].ewm(span=32).mean()
data_df['ema_100'] = data_df['close'].ewm(span=100).mean()
data_df['ema_200'] = data_df['close'].ewm(span=200).mean()
```

✓ Creating Candle Volatility Against Price Close (high - low / close)

```
data_df['vol_close'] = (data_df['high'] - data_df['low']) / data_df['close']
```

✓ Creating EMA for Candle Volatility/Price

```
# EMA for 2 hours (candles)
data_df['vol_close_ema_3'] = data_df['vol_close'].ewm(span=3, adjust=False).mean()

# EMA for 4 hours (candles)
data_df['vol_close_ema_6'] = data_df['vol_close'].ewm(span=6, adjust=False).mean()

# EMA for 8 hours ((candles)
data_df['vol_close_ema_12'] = data_df['vol_close'].ewm(span=12, adjust=False).mean()

data_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 27812 entries, 343607 to 371418
Data columns (total 20 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   date             27812 non-null   object 
 1   symbol            27812 non-null   object 
 2   open              27812 non-null   float64
 3   high              27812 non-null   float64
 4   low               27812 non-null   float64
 5   close              27812 non-null   float64
 6   ema_13             27812 non-null   float64
 7   ema_25             27812 non-null   float64
 8   ema_32             27812 non-null   float64
 9   ema_100            27812 non-null   float64
 10  ema_200            27812 non-null   float64
 11  vol_close          27812 non-null   float64
 12  vol_close_ema_3    27812 non-null   float64
 13  vol_close_ema_6    27812 non-null   float64
 14  vol_close_ema_12   27812 non-null   float64
 15  volume            27812 non-null   int64  
 16  adj_close          27812 non-null   float64
 17  adj_high            27812 non-null   float64
 18  adj_low             27812 non-null   float64
 19  adj_low_ema_3       27812 non-null   float64
```

```
6   volume usdt      27812 non-null  int64
7   tradecount      27812 non-null  int64
8   token           27812 non-null  object
9   hour            27812 non-null  int64
10  day             27812 non-null  int64
11  ema_13          27812 non-null  float64
12  ema_25          27812 non-null  float64
13  ema_32          27812 non-null  float64
14  ema_100         27812 non-null  float64
15  ema_200         27812 non-null  float64
16  vol_close       27812 non-null  float64
17  vol_close_ema_3 27812 non-null  float64
18  vol_close_ema_6 27812 non-null  float64
19  vol_close_ema_12 27812 non-null  float64
dtypes: float64(13), int64(4), object(3)
memory usage: 4.5+ MB
```

```
data_df['date'] = pd.to_datetime(data_df['date'])
```

```
columns_to_drop = ['symbol', 'token']
data_df = data_df.drop(columns=columns_to_drop)
```

```
#data_df = data_df.astype(float)
#data_df.set_index('date', inplace=True)
```

```
data_df.to_csv('data_featured.csv', index=False)
```

```
train_df = data_df.copy(deep=True)
train_df.reset_index(drop=True, inplace=True)
```

```
train_df.head()
```

	date	open	high	low	close	volume usdt	tradecount	hour	day	ema_13	ema_25	ema
0	2020-08-17 04:00:00	11844.72	11858.91	11802.35	11809.38	17196539	28800	4	1	11809.380000	11809.380000	11809.380
1	2020-08-17 05:00:00	11809.39	11836.90	11790.00	11800.01	19958274	28571	5	1	11804.334615	11804.507600	11804.548

```
#df_test = data_df[10000:]
df_test = df_test.reset_index(drop=True)
```

▼ DRL Environment Architecture

Environment Controls:

- **data**: A DataFrame containing the financial data. (passed when called the function 'env' CryptoTradingEnv(data_df))
- **take_profit_range**: A tuple representing the range for take profit levels.
- **stop_loss_range**: A tuple representing the range for stop loss levels.

- **max_stop_loss**: The maximum allowable stop loss.
- **position_size**: The maximum allowable position size.
- **initial_balance**: The initial capital.
- **balance**: Current balance.
- **position**: The current position (0 for no position, 1 for long position).
- **position_open**: The price at which the position was opened.
- **num_trades**: The number of trades.
- **profit_loss**: The cumulative profit/loss.
- **position_size**: The calculated position size as a percentage of the current balance (controlled within the code).

Observations:

- open, high, low, close, volume_usdt, hour, day, ema_13, ema_25, ema_32, ema_100, ema_200, vol_close_ema_3, vol_close_ema_6, vol_close_ema_12

Initial Reward Calculation:

- When an action is taken, the reward variable is initialized to 0.
- The variable `trade_outcome` is used to calculate the profit or loss associated with the most recent trade. It is initialized to 0.

Opening and Closing Trades:

- If the agent takes the action to open a position (action code 1), and it is not already in a trade (`self.position == 0`), it calculates the `position_size` based on the balance or maximum position size, and records the opening price.
- If the agent takes the action to close a position (action code 2), and it is currently in a trade (`self.position == 1`), it calculates the trade outcome as the difference between the opening price and the current market price.

- The profit_loss is updated with the trade outcome, representing the cumulative profit or loss across all trades.
- The balance is adjusted based on the trade outcome.

Take Profit and Stop Loss:

- The code checks whether the trade outcome is greater than 0 (a winning trade) and whether it reaches the take profit level defined by self.take_profit_range.
- If the trade reaches the take profit level (>0), a reward of 0.5 is assigned to indicate a successful trade.
- If the trade is closed within the take profit range, an extra reward of 1 is assigned to indicate a successful trade.
- If the trade outcome is less than 0 (a losing trade), the code checks for two conditions:
- If the trade outcome is equal or exceeds the maximum allowable stop loss (self.max_stop_loss * self.balance), a malus reward of -0.7 is assigned, indicating that the trade reached the maximum stop loss.
- If the trade outcome falls within the stop loss range defined by self.stop_loss_range, a reward of -1 is assigned, indicating that the trade reached the stop loss.

Trade Outcome Within Range:

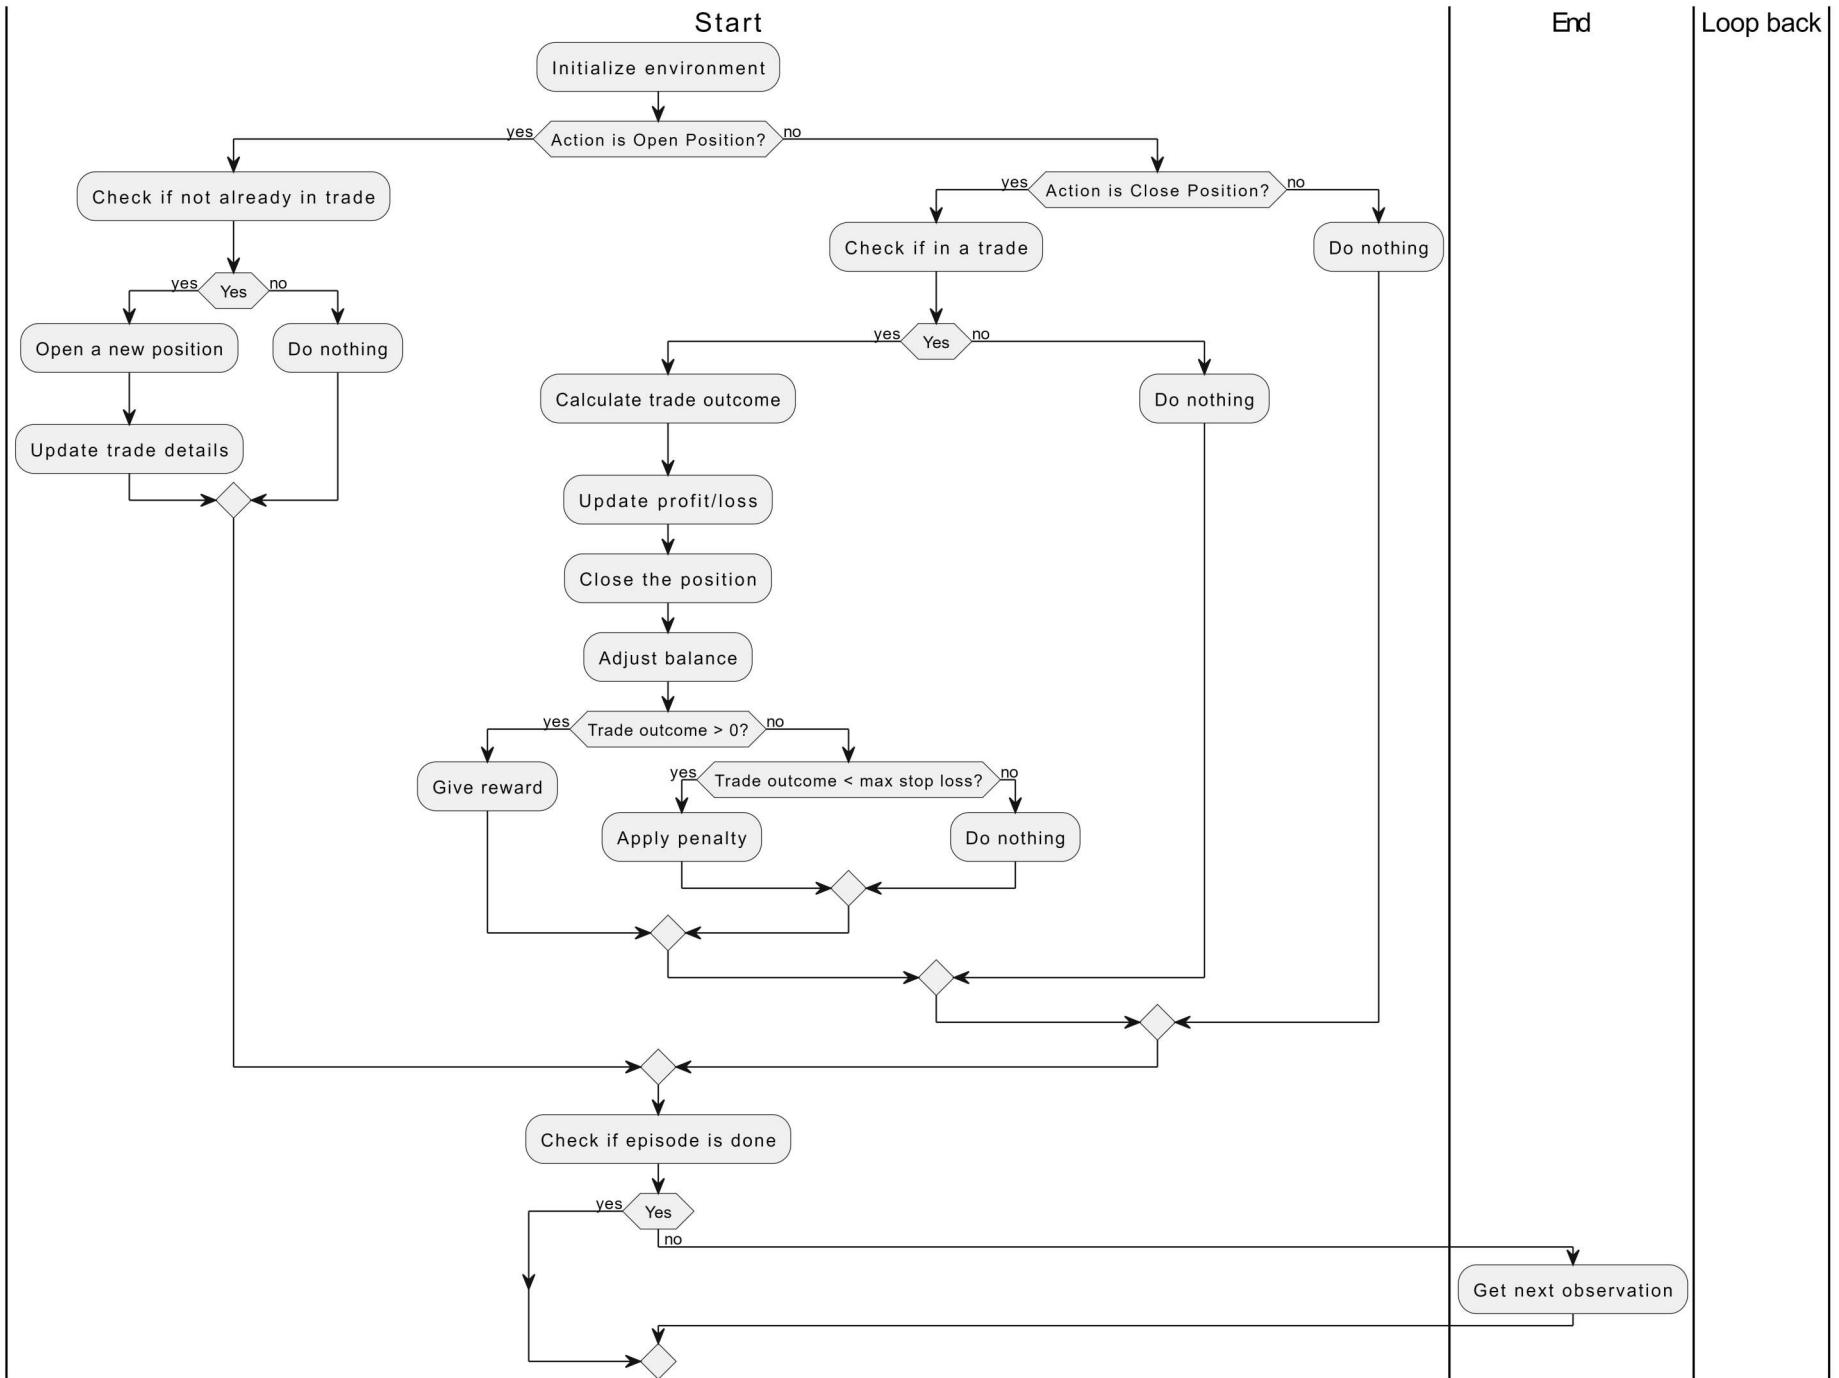
- If trade_within_range is True, it means that the trade outcome fell within the specified take profit range. In this case, self.winning_trade_within_range is incremented to keep track of such trades.

Cumulative Returns and Episode Tracking:

- The reward is also added to the episode_returns list to track rewards for the current episode.
- The cumulative_returns is updated with the reward to keep track of cumulative returns.

Completing an Episode:

- If the current step reaches the end of the trading data (`self.current_step >= self.n_steps`), the done flag is set to True, indicating that the episode has ended.



- ❖ DRL Environment v1

- It simulates cryptocurrency trading

```

class CryptoTradingEnv(gym.Env):
    def __init__(self, data, take_profit_position_range=(0.10, 0.80), stop_loss_position_range=(0.00, 0.15), max_stop_loss=0.15):
        super(CryptoTradingEnv, self).__init__()

        self.data = data
        self.n_steps = len(data)
        self.current_step = 0
        self.initial_balance = 10000
        self.balance = self.initial_balance
        self.position = 0
        self.position_open = 0 # Price at which the position was opened
        self.num_trades = 0 # Number of trades
        self.profit_loss = 0 # PnL Profit/Loss
        self.max_stop_loss_position = max_stop_loss_position

        self.take_profit_position_range = take_profit_position_range
        self.stop_loss_position_range = stop_loss_position_range

        # 0=Hold, 1=Open Position, 2=Close Position
        self.action_space = gym.spaces.Discrete(3)

        # Observations - open, high, low, and close prices, trading volume, time indicators (hour and day), and various market features
        n_features = 15
        self.observation_space = gym.spaces.Box(low=0, high=1, shape=(n_features,))

        self.episode_returns = []
        self.cumulative_returns = 0
        self.winning_trades = 0
        self.losing_trades = 0
        self.hourly_returns = []

        self.overall_rewards = 0

    def step(self, action): # This function is responsible for executing an action in the environment and returning the next state, reward, done status, and info
        self.current_step += 1
        done = False

        reward = 0

```



```

trade_outcome = 0
trade_within_range = False

if action == 1: # Open Position - checks if there is no existing position and opens a new position with a fixed p
    if self.position == 0: # Only open a position if not already in a trade
        position_size = 0.05 * self.balance
        self.position_open = self.data.loc[self.current_step, 'open']
        self.position = 1
        self.num_trades += 1
        print(f"Opened trade at step {self.current_step} with position size: {position_size:.2f}")

elif action == 2: # Close Position - checks if there is an existing position, calculates the trade outcome, update
    if self.position == 1: # Only close a position if currently in a trade
        position_close = self.data.loc[self.current_step, 'open']
        trade_outcome = position_close - self.position_open
        self.profit_loss += trade_outcome
        self.position = 0
        self.position_open = 0
        print(f"Closed trade at step {self.current_step}")
        print(f"-----> Trade outcome: {trade_outcome}")

    self.balance += trade_outcome

# Check for take profit and stop loss
if trade_outcome > 0:
    if trade_outcome >= self.take_profit_position_range[0] * self.position:
        reward = 1 # Trade reached take profit
        self.winning_trades += 1 # +1 winning trades count
        if trade_outcome <= self.take_profit_position_range[1] * self.position and trade_outcome >= self.t
            # self.trade_within_range = True
            print("Trade reached take profit")
    elif trade_outcome < 0:
        if abs(trade_outcome) >= self.max_stop_loss_position * self.position:
            reward = -1.7 # Trade reached the maximum stop loss
            self.losing_trades += 1 # +1 losing trades count
            print("Trade reached the maximum stop loss")
        elif abs(trade_outcome) >= self.stop_loss_position_range[0] * self.position:
            reward = -1 # Trade reached the stop loss
            self.losing_trades += 1 # +1 losing trades count

```



```
    print("Trade reached stop loss")
else:
```